ASK.LEGAL - LEGAL ANALYSIS EVALUATION

England & Wales Legal Analysis Model Comparison

Models	Jurisdiction	Evaluation
GPT-5.4 vs GPT-5.6 Terra	England & Wales	15 questions per arm

Scope and tested configurations

This report evaluates the 15 saved England & Wales Legal Analysis answers in each arm using **rubrics/legal-analysis-v1.md**. Each answer was scored independently before matching answers were compared side by side. To reduce ordering bias, Arm B was inspected first for odd-numbered question IDs and Arm A first for even-numbered IDs.

- **Arm A - GPT-5.4:** **gpt-5.4**, low reasoning.
- **Arm B - GPT-5.6 Terra:** **gpt-5.6-terra**, low reasoning.

The recorded configurations are otherwise identical: routing prompt, embedding model, Pinecone index and namespace, retrieval top-k, no reranker, England & Wales v2 answer prompt, Azure endpoint and API version. Both runs saved 15 successful answers. This evaluation concerns the saved answers and does not independently verify deployment state.

Independent mean scores

Scores use the rubric's 1-5 scale. Overall is the mean of all five dimensions across all 15 answers in that arm.

Arm	Legal correctness	Issue coverage	Grounding & citations	Jurisdiction discipline	Usefulness & safety	Overall
Arm A - GPT-5.4	3.87	4.13	3.00	5.00	4.20	4.04
Arm B - GPT-5.6 Terra	4.47	4.67	4.13	4.93	5.00	4.64

The largest mean difference was grounding and citations, where Arm B scored 4.13 and Arm A scored 3.00. Both arms remained strongly disciplined to England and Wales.

Pairwise preference totals

Preferred arm	Questions
Arm A	1
Arm B	13
Tie	1
Total	**15**

Of Arm B's 13 preferences, ten were confident and three were close. Arm A's single preference was close; the tie was close.

Question-level comparisons

Question	Preference	Confidence	Key reason
ew-001	Arm B	Confident	Adds the distinct ERA s 104 statutory-right dismissal route and a fuller alternative-dismissal analysis.
ew-002	Arm A	Close	Both omit core 1985 provisions, but Arm A includes deposit non-protection, prescribed information and the one-to-three-times sanction.
ew-003	Arm B	Confident	Gives the clearer CRA remedy sequence and expressly explains the first-six-month presumption.
ew-004	Arm B	Confident	Reaches the same contract/damages result with cleaner authority and no internal-source citations.
ew-005	Arm B	Confident	Leads safely with the imminent limitation issue; Arm A contains an impossible future-dated citation.
ew-006	Arm B	Confident	Better separates child, CMS, spousal, housing, pension and overall-settlement regimes.
ew-007	Tie	Close	Arm A catches domicile; Arm B states the cohabitant maintenance standard. Each omits one material threshold point.
ew-008	Arm B	Confident	Anchors conflicts, unfair prejudice, derivative standing and reflective loss in the correct statutory routes.
ew-009	Arm B	Close	Arm B is more proportionate and operationally safe; Arm A spots more special-category issues but relies on dated guidance.
ew-010	Arm B	Close	Both are strong, but Arm B is cleaner, more concise and more cautious about litigation and injunctions.
ew-011	Arm B	Confident	States the right of a voluntary attendee to leave and distinguishes bringing, surrendering, unlocking and extracting the phone.
ew-012	Arm B	Confident	Focuses on proving the French will and obtaining the correct English grant rather than peripheral administration jurisdiction.
ew-013	Arm B	Confident	Adds the debtor-status-dependent Debt Claims Protocol and a better warning against insolvency pressure.
ew-014	Arm B	Close	Adds insurance-protection and limitation cautions and avoids an oversimplified court-forum threshold.
ew-015	Arm B	Confident	Gives a more commercially useful German asset-preservation analysis with cleaner authority.

Recurring strengths and weaknesses

Arm A - GPT-5.4

Strengths:

- Comprehensive issue spotting and detailed evidence checklists.
- Consistent England-and-Wales jurisdiction discipline.
- Frequent attention to causation, mitigation, limitation and enforcement practicality.

Weaknesses:

- Long, sometimes dated or peripheral authority lists often obscure the governing rule.
- Internal retrieval filenames are presented as authority in ew-004, ew-010 and ew-015.
- Key routes are omitted or underdeveloped in several answers, including ERA s 104, CA 2006 ss 994/996, the voluntary attendee's right to leave and the Debt Claims Protocol.

- Some procedural guidance is incomplete or overbroad, including the court-forum statement in ew-014 and cross-border choice-of-law treatment in ew-012.

Arm B - GPT-5.6 Terra

Strengths:

- Clear separation of claims, parties, forums and remedies.
- More direct use of governing statutory provisions and leading authorities.
- Concise, safe action plans that foreground urgent deadlines and evidence preservation.
- Strong treatment of claim ownership and procedural posture in company, criminal interview and cross-border disputes.

Weaknesses:

- Some important secondary frameworks remain omitted: the principal 1985 housing duties in ew-002, the domicile threshold in ew-007, DPIA specifics in ew-009 and Hague 2005 recognition in ew-015.
- The cross-border succession answer does not fully distinguish administration law from the law governing beneficial succession.
- Several recent citations could not be verified within the supplied record.

Critical failures

- **Arm A:** One. The ew-005 reference list cites **Ringe v Eden Springs (UK) Ltd [2112] EWHC 0014 (QB)**, an impossible future-dated citation. The substantive limitation advice in that answer remains usable, but the citation is a critical grounding failure.
- **Arm B:** None identified.

Limitations

- No expectation file or gold answers were available.
- No retrieval logs or gateway logs were available, so citation support and retrieved context could not be checked.
- No external source checking was performed. A plausible-looking citation was not treated as proof of accuracy, and recent authorities remain unverified unless an error was internally apparent.
- One **gpt-5.6-sol** LLM judge evaluated this jurisdiction and saw the arm identifiers and model names. Model identity was therefore not blinded, although answer-inspection order was alternated; there was no second judge or inter-rater agreement check.
- There was one saved run per arm and no repeated runs, so response variability was not assessed.
- The 15-question suite is purposive and does not support statistical inference.
- The evaluation does not independently verify which recorded backend configuration was deployed.

These are findings about the saved answers only. No deployment or promotion recommendation is made.