ASK.LEGAL - LEGAL ANALYSIS EVALUATION

Hong Kong Legal Analysis Model Comparison

Models	Jurisdiction	Evaluation
GPT-5.4 vs GPT-5.6 Terra	Hong Kong	15 questions per arm

1. Scope and tested configurations

This is an independent LLM-as-judge evaluation of 15 saved Hong Kong Ask.Legal legal-analysis answers per arm. Each answer was scored independently under rubrics/legal-analysis-v1.md before the matching pair was compared. To reduce ordering bias, arm-b was inspected first for odd-numbered question IDs and arm-a first for even-numbered IDs.

The experiment compares **GPT-5.4** (arm-a) with **GPT-5.6 Terra** (arm-b). Both recorded configurations use **low reasoning effort** and are otherwise identical: intent-router-current, text-embedding-3-small, ask-legal-chatbot-hk-dataset-v8 in the default namespace, retrieval top-k 8, no reranker, and hong-kong-v2.0 as the answer prompt. The recorded main **model** field is the only differing configuration field.

This report evaluates saved answers, not legal advice. It does not independently establish that the recorded configurations were deployed.

2. Independent scores

Means are arithmetic means across all 15 answers for each arm. Overall is the mean of all five dimensions across the 15 answers.

Arm	Legal correctness	Issue coverage	Grounding and citations	Jurisdiction discipline	Usefulness and safety	Overall
arm-a - GPT-5.4	4.07	4.60	3.60	4.93	4.53	4.35
arm-b - GPT-5.6 Terra	4.47	4.80	3.93	4.80	4.73	4.55

3. Pairwise preference totals

Preference	Count
arm-a - GPT-5.4	2
arm-b - GPT-5.6 Terra	11
Tie	2
Total	**15**

4. Question-level comparisons

Question	Preference	Confidence	Key reason
hk-001	arm-b	Close	More accurately confines statutory retaliation protection while separating ordinary commission and notice claims.
hk-002	arm-b	Close	Reaches the same sound result with tighter causation analysis and fewer weakly connected authorities.
hk-003	arm-b	Close	Better explains the post-acceptance Sale of Goods Ordinance s 55 remedy and avoids unsupported additional statutes.
hk-004	arm-b	Close	More direct on electronic formation and the value credit due for any usable partial work.
hk-005	Tie	Close	Both give sound negligence and limitation guidance but wrongly route low-value personal-injury claims to Small Claims.
hk-006	arm-b	Close	More completely separates parenting arrangements, child maintenance and spousal maintenance.
hk-007	arm-a	Confident	Correctly checks domicile before applying succession law to movable bank balances; arm-b does not.
hk-008	Tie	Close	Both give strong minority-remedy analysis but wrongly say an exact 25% holding can block a 75% special resolution.
hk-009	arm-b	Close	More distinctly addresses later use, retention and security of monitoring data.
hk-010	arm-b	Confident	Gives a materially safer account of the interim-injunction threshold and fuller privilege and malicious-falsehood analysis.
hk-011	arm-b	Close	Ties safe interview and lawyer-access advice more directly to the questioning directions.
hk-012	arm-b	Close	More focused on foreign-domicile evidence and the estate-duty cutoff.
hk-013	arm-a	Close	Uniquely checks arbitration and forum clauses before selecting the Small Claims Tribunal.
hk-014	arm-b	Confident	Gives the better loss analysis and avoids arm-a's dated direction to HKICPA, although it should still name the current regulator.
hk-015	arm-b	Confident	Makes Mainland asset preservation and reciprocal enforcement central to a Shenzhen-assets dispute.

5. Recurring strengths and weaknesses

Arm-a - GPT-5.4

Strengths:

- Thorough issue spotting, evidence-preservation checklists and practical action plans.
- Strong conflict-of-laws discipline on hk-007.
- Useful attention to arbitration and forum clauses on hk-013.
- Consistently stays within Hong Kong law and usually qualifies conclusions appropriately.

Weaknesses:

- Frequently includes long authority lists whose connection to the actual proposition or facts is weak, reducing auditability.
- Incorrectly places low-value personal-injury claims in the Small Claims Tribunal on hk-005.
- Incorrectly treats exactly 25% of voting rights as enough to block a 75%-threshold special resolution on hk-008.
- Gives a dated accounting-complaint route on hk-014 and is occasionally insufficiently cautious about exceptional interim relief.

Arm-b - GPT-5.6 Terra

Strengths:

- More selective and concise use of legal rules and authorities.
- Strong distinctions between available remedies and procedural risks, especially on consumer rights, defamation, police interviews and cross-border enforcement.
- Practical next steps are consistently tied to the relevant facts.
- Generally stronger legal precision and issue coverage, reflected in the independent means.

Weaknesses:

- Shares arm-a's incorrect Small Claims guidance on hk-005.
- Shares arm-a's incorrect special-resolution blocking advice for an exact 25% holding on hk-008.
- Fails to check domicile before applying intestacy rules to the bank accounts on hk-007.
- Does not name or adequately ground the current accounting regulator on hk-014.

6. Critical failures

Each arm has one critical failure:

- **arm-a, hk-005:** incorrectly states that a personal-injury claim under HK$75,000 may be brought in the Small Claims Tribunal, which does not have jurisdiction over personal-injury claims.
- **arm-b, hk-005:** the same incorrect Small Claims Tribunal guidance.

The hk-008 special-resolution error is also material and appears in both arms, but it was recorded as a substantial correctness defect rather than an additional critical failure.

7. Limitations

- No expectation file or gold answers were available.
- No retrieval logs or sanitized gateway logs were available.
- No external source checking was performed. Citations were not treated as accurate merely because they appeared plausible, and their existence, currency and proposition-level support could not be independently verified.
- One **gpt-5.6-sol** LLM judge evaluated this jurisdiction and saw the arm identifiers and model names. Model identity was therefore not blinded, although answer-inspection order was alternated; there was no second judge or inter-rater agreement check.
- Only one saved run per arm was available; there were no repeated runs.
- The 15-question suite does not support statistical inference.
- Differences cannot be attributed causally beyond the recorded model change, and the recorded deployment was not independently verified.

8. Findings only

The recorded results show a higher independent overall mean for arm-b, 4.55 versus 4.35, and pairwise preferences of 11 for arm-b, 2 for arm-a and 2 ties. Arm-b's advantage is primarily in precision, selective grounding and remedy analysis. Arm-a remains materially stronger on the hk-007 conflict-of-laws issue and narrowly stronger on hk-013's dispute-resolution caveat. Both arms require correction of the shared hk-005 and hk-008 errors.

This report stops at comparative findings and makes no deployment, promotion or release recommendation.