ASK.LEGAL - LEGAL ANALYSIS EVALUATION

Singapore Legal Analysis Model Comparison

Models	Jurisdiction	Evaluation
GPT-5.4 vs GPT-5.6 Terra	Singapore	15 questions per arm

Scope and tested configurations

This saved-answer evaluation compares **GPT-5.4 (arm-a)** with **GPT-5.6 Terra (arm-b)** on 15 Singapore Legal Analysis questions. Both arms recorded **low reasoning effort** and otherwise identical configuration: the same routing prompt, embedding model, Singapore Pinecone index and namespace, retrieval top-k, no reranker, Singapore v1.4 answer prompt, Azure endpoint, and API version. The recorded main model is the only configuration difference.

Each answer was scored independently under **rubrics/legal-analysis-v1.md** before matching answers were compared side by side. This is an evaluation of saved answers, not legal advice.

Independent scores

Arm	Legal correctness	Issue coverage	Grounding and citations	Jurisdiction discipline	Usefulness and safety	Overall
arm-a - GPT-5.4	4.07	4.33	3.87	4.80	4.40	4.29
arm-b - GPT-5.6 Terra	4.20	4.67	3.73	4.93	4.60	4.43

Means are arithmetic means across all 15 independently scored answers for each arm. Overall is the mean of all five dimensions across those answers.

Pairwise preference totals

Preference	Questions
arm-a	4
arm-b	8
Tie	3
Total	**15**

Question-level comparisons

Question	Preference	Confidence	Key reason
sg-001	arm-b	confident	Gives the one-month wrongful-dismissal and usual six-month post-employment salary deadlines; arm-a gives only generic urgency.
sg-002	tie	close	Arm-a is better grounded; arm-b is more complete on evidence, losses, itemisation, and possible SCT use.
sg-003	tie	close	Both accurately apply the CPFTA six-month presumption and full repair, replacement, price-reduction, and rescission sequence.
sg-004	arm-b	close	Better coverage of repudiation, net completion costs, IP/work product, system security, and limitation.
sg-005	tie	close	Both are strong on negligence, causation, contributory fault, damages, evidence, and limitation, but omit concrete pre-action protocol steps.
sg-006	arm-b	close	More complete on HDB transfer/buy-out, CPF allocation, parenting, and package settlement; both have current-procedure and maintenance-eligibility defects.
sg-007	arm-b	confident	Identifies Muslim status as a potentially dispositive fact and includes the administration process; arm-a omits that regime check.
sg-008	arm-a	confident	Similar substantive quality, but arm-a grounds oppression and derivative relief in statute and cases; arm-b cites none.
sg-009	arm-b	close	More focused and avoids arm-a's questionable six-year retention suggestion; both overstate mutual trust and confidence.
sg-010	arm-a	confident	Avoids arm-b's inadequately supported three-year knowledge-based defamation limitation statement.
sg-011	arm-a	confident	Explains CPC attendance and the protection against self-incriminating answers; arm-b omits both despite stronger phone-power coverage.
sg-012	arm-b	confident	Adds the direct Singapore-grant route and the current 2024 probate rule to the sound re-sealing analysis.
sg-013	arm-a	close	Adds more precise SOPA response rules, the Magistrate's Court route, and a particularly usable defects matrix.
sg-014	arm-b	close	Similar civil analysis, but arm-b more clearly distinguishes professional misconduct from negligence; both incompletely treat Noor Azlin's history.
sg-015	arm-b	close	Gives more precise s 12A conditions, freezing evidence, wrongful-termination risk, anti-suit relief, and Indonesia-focused preservation planning.

Recurring strengths and weaknesses

Arm-a - GPT-5.4

Strengths:

- Often provides dense statutory and case grounding, particularly in corporate law, criminal procedure, consumer remedies, and construction payment.
- Consistently offers detailed evidence-preservation steps and practical claim framing.
- Its strongest answers identify high-impact procedural protections, notably the police interviewee's self-incrimination protection.

Weaknesses:

- Occasionally omits a decisive fact, deadline, alternative route, or statutory power: examples include the exact employment deadline, Muslim intestacy, direct Singapore probate, and CPC phone access.

- Some procedural and case-history material is outdated or incomplete, including the Family Justice Rules 2014 references and the first-instance Noor Azlin causation disposition.
- Several answers are more diffuse than necessary for the practical decision.

Arm-b - GPT-5.6 Terra

Strengths:

- More consistently checks threshold facts and alternative routes, including Muslim intestacy, direct Singapore probate, IP/system security, and Indonesia-focused asset protection.
- Action plans are generally focused on the most urgent practical risk.
- Often distinguishes related legal routes clearly, including oppression versus derivative recovery and negligence versus professional discipline.

Weaknesses:

- Citation density is uneven: sg-008 has no authorities, while several PDPA propositions lack controlling sections.
- Sg-010 gives an inadequately supported three-year defamation limitation statement, and sg-011 omits self-incrimination and compulsory-attendance protections.
- It shares currency defects with arm-a on mutual trust and confidence, Family Justice Rules 2014, and the later Noor Azlin history.

Critical failures

No critical failures were recorded for either arm. The identified defects were reflected in dimension scores and, where material to the pair, in the comparison findings.

Limitations

- No expectation file or gold answers were available.
- No retrieval or gateway logs were available, so the retrieved basis for an authority and the accuracy of citations could not be checked against source context.
- No external source checking was performed. Citation plausibility was not treated as proof of accuracy; findings rely only on repository inputs and the judge's legal analysis.
- One **gpt-5.6-sol** LLM judge evaluated this jurisdiction and saw the arm identifiers and model names. Model identity was therefore not blinded, although answer-inspection order was alternated; there was no second judge or inter-rater agreement check.
- Each arm has one saved run only. There are no repeated runs for assessing model-response variability.
- The 15-question suite supports descriptive findings only. No statistical inference is warranted.
- The evaluation relies on the recorded configurations and does not independently verify what was deployed.

These are comparative findings only. No deployment, promotion, or release recommendation is made.